




Exploring DoRA: Weight-Decomposed Low-Rank Adaptation Across Modalities

Richie Xue • Shaurya Sen • Kyle Du


Cornell University · Dept. of Computer Science · CS4782: Introduction to Deep Learning · Spring 2025

The Problem

Fine-tuning billion-parameter models is costly for most teams. **LoRA** (Hu et al., 2022) helps by freezing the base weights W_0 and learning a low-rank update BA , but this update changes a layer's *magnitude* (how strong) and *direction* (what feature) **together**.

DoRA (Liu et al., 2024) decouples them: learn magnitude as a separate scalar, and let LoRA update only the direction.

DoRA Architecture




DoRA splits each weight W_0 into magnitude m and direction $V/\|V\|_c$, applying a LoRA update BA only to the direction while training m as a separate scalar.

Audio Results

Wav2Vec2-base (67M params) fine-tuned on Google Speech Commands v0.02 — 84.8k utterances, **12 keyword classes**.

Method	Val Acc	Test Acc	Trainable	Time
DoRA ($r=8$)	98.6%	89.7%	826.6k	183.6 min
LoRA ($r=8$)	98.5%	89.0%	789.8k	190.8 min

+0.7% test accuracy. The magnitude scalar's small overhead is offset by faster directional gradient convergence on the audio encoder.



-44.6% average train loss. DoRA shows lower average training loss, suggesting faster optimization and smoother convergence.

NLP Results


Cross-Task at $r=8$: (RoBERTa-base accuracy / F1+acc / accuracy)

Task	DoRA	LoRA	FFT	Δ
SST-2 (67k, sentiment)	93.2%	93.3%	93.2%	-0.1
MRPC (3.7k, paraphrase)	87.8%	85.9%	74.8%	+1.9
RTE (2.5k, inference)	70.0%	66.1%	54.9%	+3.9
SST-2 <i>large</i>	96.4%	95.5%	95.4%	+0.9
MRPC <i>large</i>	89.3%	89.5%	80.6%	-0.2
RTE <i>large</i>	76.9%	78.7%	68.6%	-1.8
OpenLLaMA-3B	81.0% [†]	—	—	—

DoRA's advantage is largest on **low-data tasks** (RTE +3.9 pp) where magnitude decoupling reduces directional gradient interference. On large models both methods saturate.


Benchmark Task Examples: (Text $\rightarrow$ Label)

SST-2 "...crowd-pleasing as a documentary can get." $\rightarrow$ **Positive**
MRPC S1: "Corixa rose 54c to \$7.74." S2: "...rose 54c, or 8%, to \$7.74." $\rightarrow$ **Not Equiv.**
RTE P: "Leaders signed Egypt-Israel treaty in 1979." H: "Treaty signed in 1979." $\rightarrow$ **Not Entailed**



DoRA converges faster and to a lower eval loss than LoRA on RTE, while full fine-tuning stalls near the random-guess baseline.

Magnitude-Direction Updates: (Full vs. LoRA vs. DoRA)



Full $r=0.40$, LoRA $r=0.93$, DoRA $r=0.24$. DoRA shows significantly lower slope and correlation than LoRA, closer to FFT, showing the effect of de-coupling **M** and **D**.


Vision & Robot Results

Object Grasping: Predict a 2-D bounding box around the graspable object in an RGB tabletop image. Frozen ViT-B/16 and SigLIP-B/16 backbones, adapted at $r=8$, scored by IoU.

Backbone	DoRA	LoRA
ViT-B/16	11.9% IoU	11.9% IoU
SigLIP-B/16	22.0% IoU	21.5% IoU

SigLIP's richer vision-language pretraining benefits more from magnitude decoupling. ViT features show minimal difference.

Qualitative: SigLIP + DoRA: Green = predicted grasp Red = ground truth



Push-T VLA Task: A robot slides a T-shaped block to a target pose from an overhead camera plus language instruction. DoRA adapts a 126-layer OpenVLA model; action MSE drops 63.6 $\rightarrow$ **27.9** over 3 epochs.



DoRA adapter driving Push-T rollouts

Conclusions

- **Low-data NLP:** DoRA gains most when data is scarce — RTE +3.9 pp, MRPC +3.0 pp at $r=16$.
- **Audio:** +0.7% test acc. and 25% faster training on Wav2Vec2 — directional convergence is faster.
- **Vision:** SigLIP amplifies DoRA's benefit (+0.5 pp IoU); ViT-base does not, meaning base models matter.

Future Work

- **QLoRA + DoRA:** 4-bit quantized fine-tuning to reach 7B+ models on consumer hardware.
- **Diffusion models:** apply DoRA to action-distribution heads and image-generation U-Nets.
- **SVD-based decomposition:** explore EDoRA [7] — uses SVD to decouple singular directions, potentially reducing the parameter budget beyond DoRA.