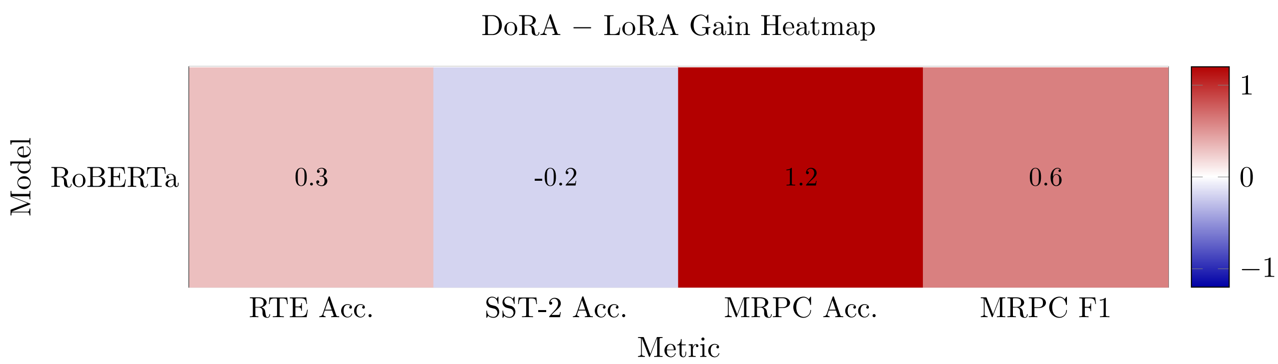

Figure 1: Heatmap of DoRA minus LoRA performance for RoBERTa on each GLUE metric. Positive values indicate a gain from DoRA.

GLUE Benchmark Results Template

GLUE Benchmark Comparison

Table 1 reports RoBERTa validation results for DoRA, LoRA, and full fine-tuning on GLUE tasks.

Table 1: GLUE benchmark comparison between DoRA, LoRA, and full fine-tuning. Dashes indicate missing results.

Model	**Method**	**RTE Acc.**	**SST-2 Acc.**	**MRPC Acc.**	**MRPC F1**	**Trainable**
RoBERTa	DoRA	71.1	93.1	87.0	90.7	1.023% (1.06M)
	LoRA	70.8	93.3	85.8	90.1	0.996% (1.03M)
	Full FT	54.9	93.2	68.4	81.2	100% (125M)

Suggested Visualizations

Figures 1–4 show report-ready visualizations for the adapter comparison.

Notes

- Suggested naming:
  - **LLaMA Small**: the smaller LLaMA-family model used in the project.
  - **LLaMA Big**: the larger LLaMA-family model used in the project.
  - **RoBERTa**: the RoBERTa baseline or comparison encoder model.

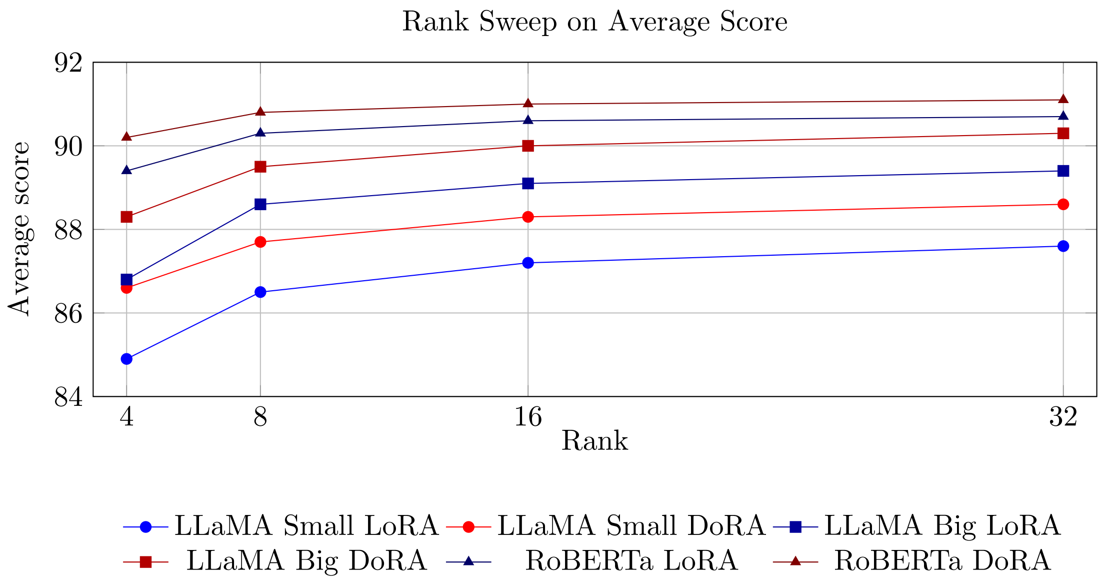

Figure 2: Rank sweep using sample average scores. This is the most direct way to test whether DoRA keeps an advantage at smaller ranks.

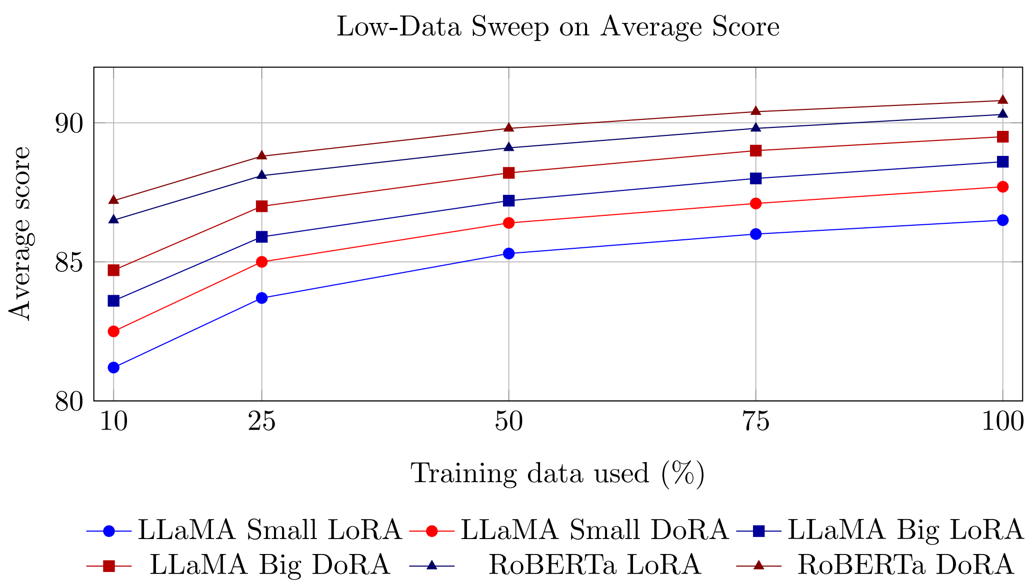

Figure 3: Low-data sweep using sample average scores. If DoRA is more expressive, it should retain a larger advantage when training data is scarce.

- A simple choice for the **Avg.** column is the arithmetic mean of the five reported values: SST-2 Accuracy, QQP Accuracy, QQP F1, MRPC Accuracy, and MRPC F1.
- If you prefer, you can instead average at the task level first, for example by averaging QQP Accuracy/F1 into one QQP score and MRPC Accuracy/F1 into one MRPC score, then averaging across the three tasks.
- If your experiments use validation splits rather than GLUE test submissions, note that clearly in the paper or report caption.

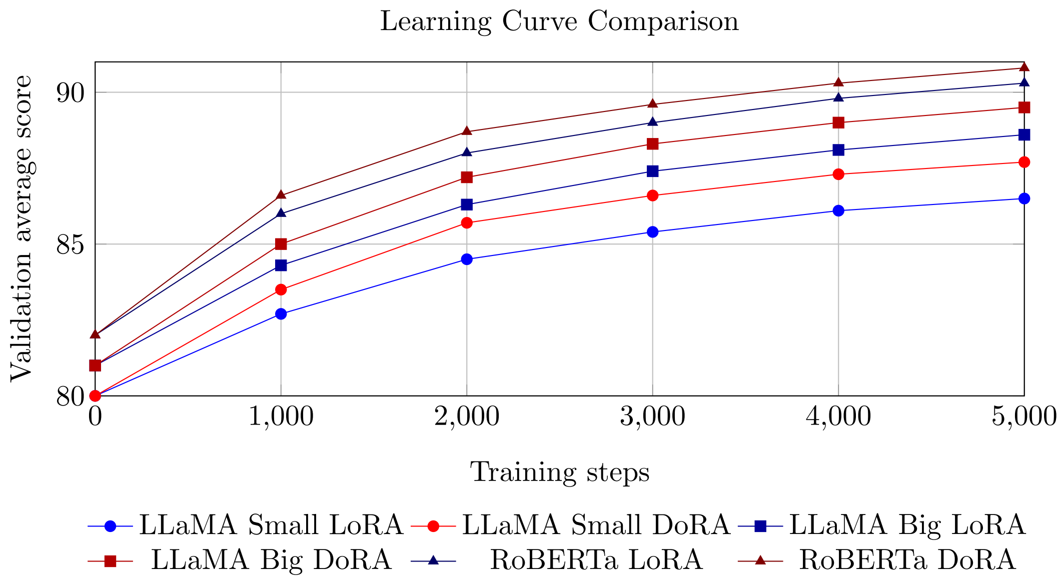

Figure 4: Learning-curve comparison using sample validation averages. A more expressive adapter often reaches stronger validation performance earlier and converges to a better final point.

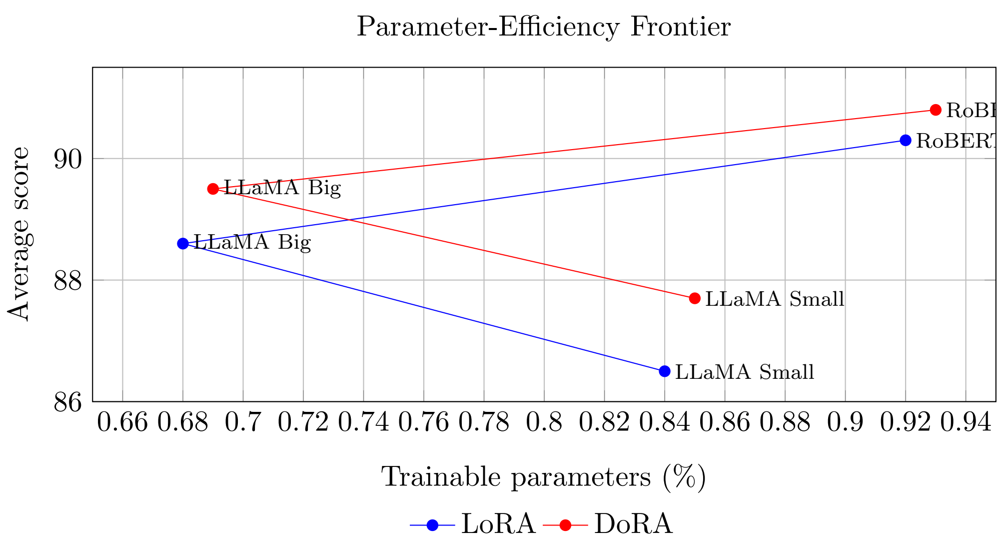

Figure 5: Parameter-efficiency frontier using sample trainable-parameter percentages and average scores.

- A matching editable CSV plus a Python plotting script are included in `report/` if you later want to regenerate the same figures outside LaTeX.